

Machine Learning 1 – Fundamentals

Decision Trees

Prof. Dr. J. M. Zöllner, M.Sc. Nikolai Polley, M.Sc. Marcus Fechner

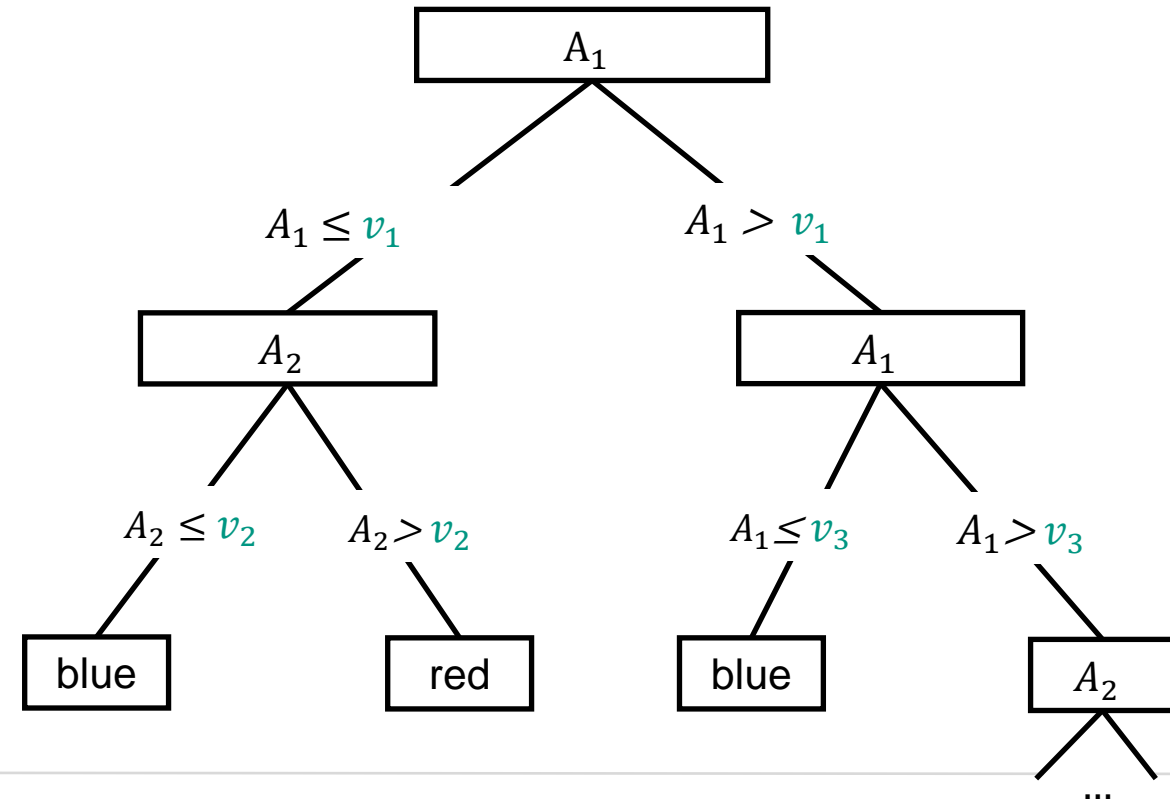



Overview

- Motivation
- Formalization
- Attribute selection
- Build the Tree - ID3-algorithm
- Overfitting
- Bagging
- Random Forest
- Extensions

Motivation

- **Classification:** Partition the space so that all instances in a region have the same class
- Extensible to regression trees

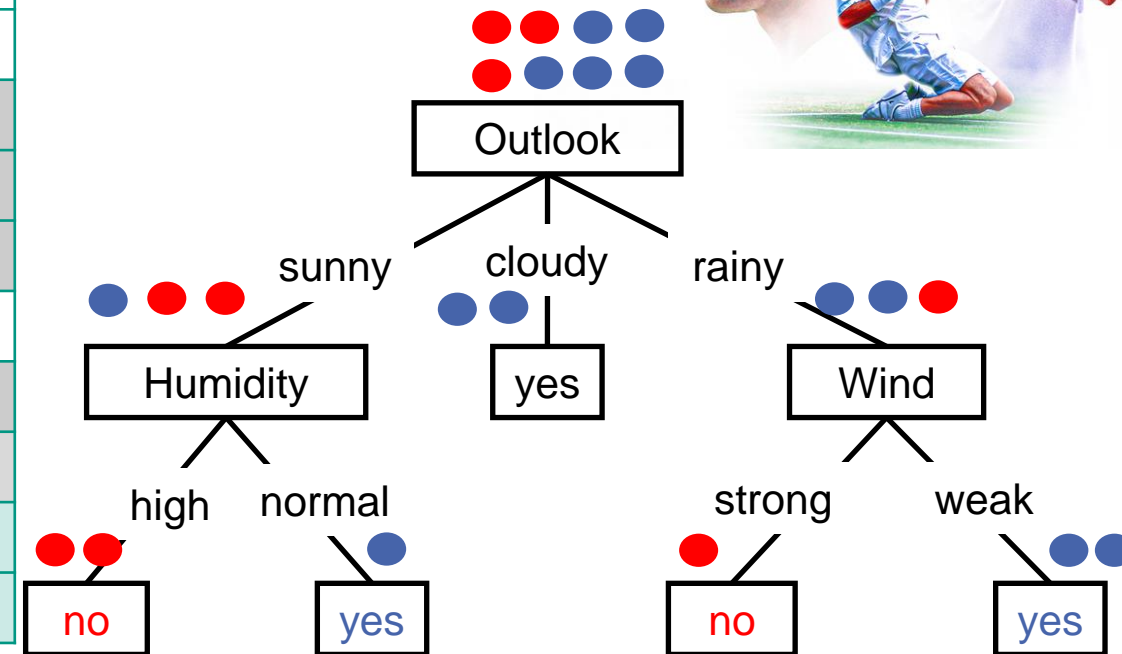


Motivation: Play Tennis

■ **Question:** Which days are suitable for Roger Federer to play tennis?



No	Outlook	Temperature	Humidity	Wind	Tennis?
1	Sunny	Hot	High	Weak	No
2	Sunny	Hot	High	Strong	No
3	Cloudy	Hot	High	Weak	Yes
4	Rainy	Warm	High	Weak	Yes
5	Rainy	Cold	Normal	Weak	Yes
6	Rainy	Cold	Normal	Strong	No
7	Cloudy	Cold	Normal	Strong	Yes
8	Sunny	Cold	Normal	Weak	Yes
9	Sunny	Warm	High	Weak	???
10	Rainy	Warm	Normal	Weak	???

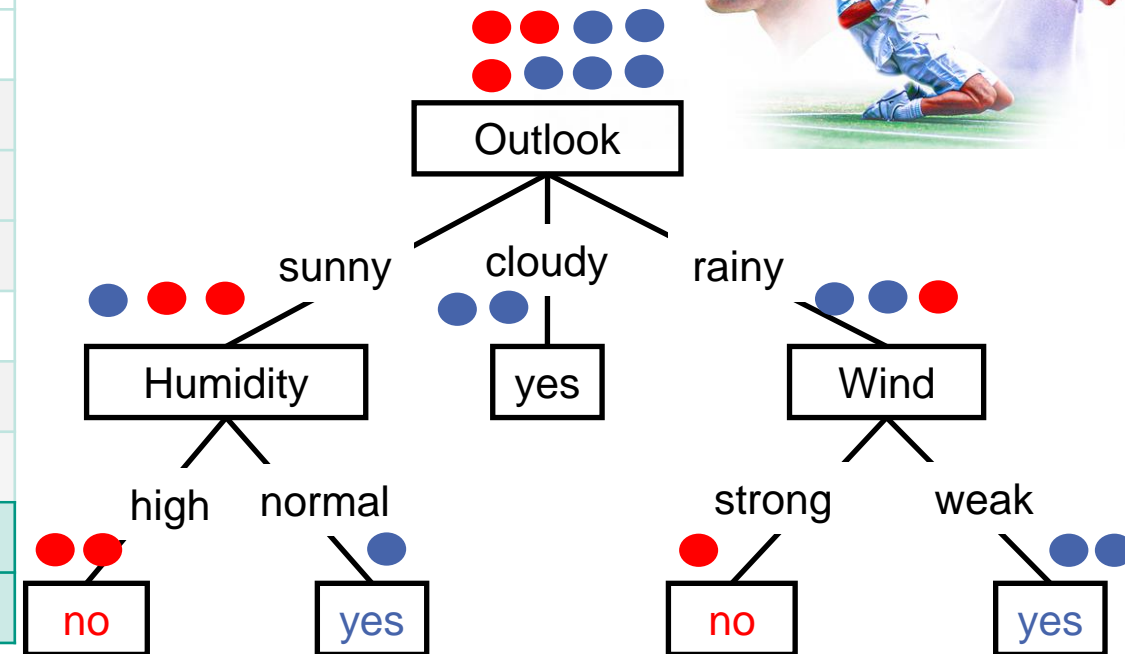


Motivation: Play Tennis

■ **Question:** Which days are suitable for Roger Federer to play tennis?



No	Outlook	Temperature	Humidity	Wind	Tennis?
1	Sunny	Hot	High	Weak	No
2	Sunny	Hot	High	Strong	No
3	Cloudy	Hot	High	Weak	Yes
4	Rainy	Warm	High	Weak	Yes
5	Rainy	Cold	Normal	Weak	Yes
6	Rainy	Cold	Normal	Strong	No
7	Cloudy	Cold	Normal	Strong	Yes
8	Sunny	Cold	Normal	Weak	Yes
9	Sunny	Warm	High	Weak	No
10	Rainy	Warm	Normal	Weak	Yes



When are Decision Trees suitable?

Basic Decision Trees:

- Instances are represented as attribute-value pairs (categorical)
- Target function returns discrete output values (e.g. classes)
- Non-parametric: No assumptions about underlying distribution
- Interpretability is possible
- Generally low computing resource requirements for inference
- Simple to learn non-linear problems
- Data doesn't need to be normalized and can be in diverse formats

Decision Tree with extensions:

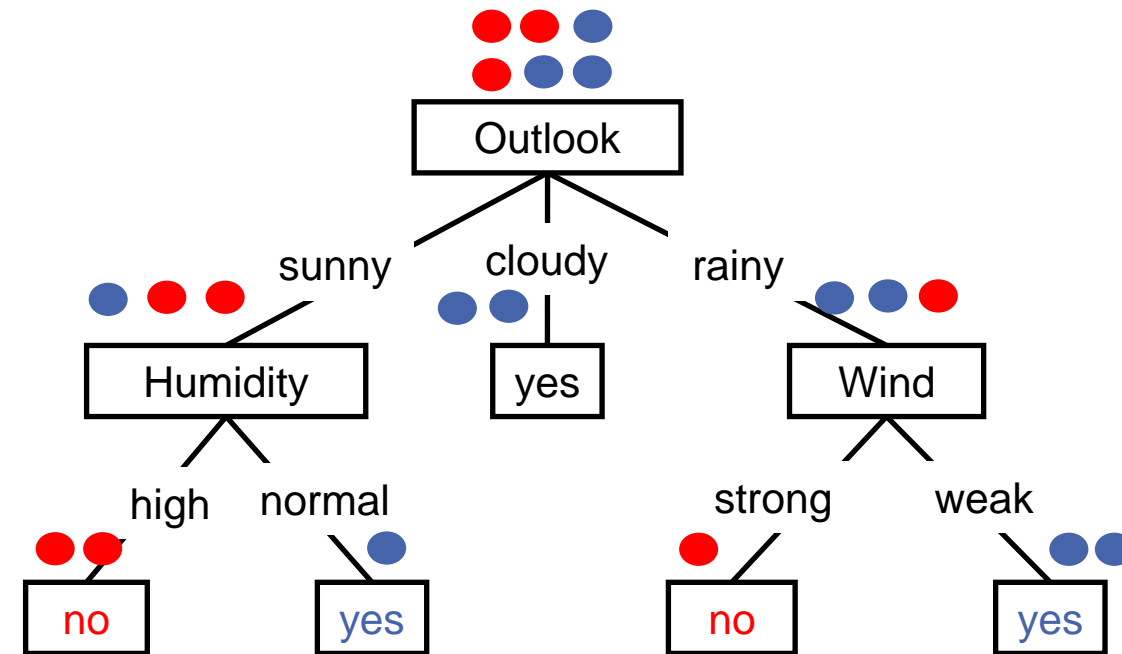
- Applicable for regression (Regression Tree)
- Noisy input data
- Continuous attribute values (features)
- Missing input data

Overview

- Motivation
- Formalization
- Attribute selection
- Build the Tree - ID3-algorithm
- Overfitting
- Bagging
- Random Forest
- Extensions

Formalization

- **Nodes:** test feature/attribute A for value
 - E.g. Outlook, Humidity, Wind ...
- **Root Node:** Top/first node in decision tree
- **Child Node:** Successor of a node
- **Parent Node:** Predecessor of a node
- **Leaf:** Final node containing classification result Y
 - E.g.: PlayTennis?=yes, PlayTennis?=no
- **Branch:** Attribute value v of testing A
 - E.g., for the attribute Outlook: sunny, cloudy, rainy



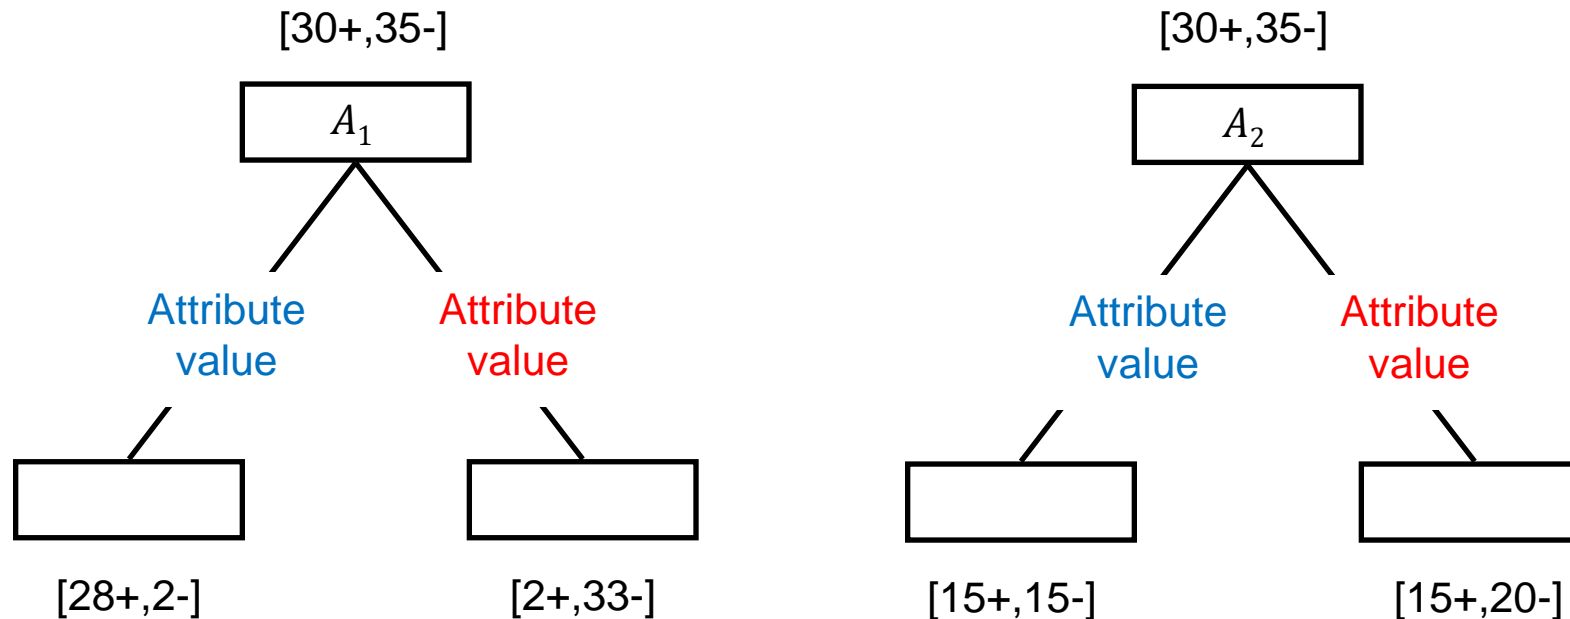
Overview

- Motivation
- Formalization
- Attribute selection
- Build the Tree - ID3-algorithm
- Overfitting
- Bagging
- Random Forest
- Extensions

Attribute Selection

- **Problem:** How can we measure how well an attribute/feature splits the data?
- **Example:** In the following, is attribute A_1 or A_2 preferable?
- **Solution:** Metrics to measure impurity of data sets

- Entropy
- Gini-Index
- ...



Notation: [number positive examples (+); number negative examples (-)]


Entropy

- **Definition:** *In information theory, the entropy of a random variable Y is the average level of "information", "surprise", or "uncertainty" inherent to the variable's possible outcomes y_i*
随机变量 Y 的熵是指变量可能结果 y_i 所固有的“信息”、“惊奇”或“不确定性”的平均水平
- **For decision trees:** Entropy is a measure of the homogeneity (in terms of class membership) of the current data S
- **Here:** Discrete random variable $Y \sim p(y)$, whereas $p(y) = \frac{|S_y|}{|S|}$ defined by data S and subset S_y of data with class y
 - Entropy for **K Classes**: $H(S) = -\sum_{i=1}^K p(y_i) \log_2 p(y_i)$
 - Entropy for **2 Classes**(\oplus, \ominus): $H(S) = -p_{\oplus} \log_2 p_{\oplus} - p_{\ominus} \log_2 p_{\ominus}$
- **Info:** The entropy, measured with a logarithm of base 2 uses **Bit** as unit

Entropy – Decision Tree



■ High entropy

- Y is derived from an almost uniformly distributed probability density
- Data sampled from this density are difficult to predict





■ Low entropy

- Y is derived from a probability density with a high probability for one class, relative to the others
- Data sampled from this density are highly predictable



Entropy – Decision Tree

- **Objective:** To quickly sort data into their respective classes through the selection of appropriate attributes
 - i.e. successively reduce the entropy as quickly as possible




Information Gain

- **Information Gain** $IG(S, A)$: Expected entropy reduction of S by splitting on attribute A
- $IG(S, A) = H(S) - \sum_{v \in V(A)} \frac{|S_v|}{|S|} H(S_v)$
 - $V(A)$: Set of all possible attribute values v of attribute A
 - S_v : Subset of S for which A is value v
 - Calculate the difference between the entropy of the parent node and the entropy of the child nodes
- **Objective of learning with decision trees**
 - Select attributes that minimize entropy and maximize information gain.
Classify learning examples with as few steps as possible
→ Tree with little depth

Information Gain – Attribute Selection

- **Idea:** A suitable attribute splits the instances into subsets, where (ideally) all instances in a subset are positive or negative



Information Gain – Attribute Selection

■ Step 1: Calculate entropy of parent node

■ Entropy of S :

$$■ H(S) = -p_{\oplus} \log_2 p_{\oplus} - p_{\ominus} \log_2 p_{\ominus}$$

$$■ H(S) = -\frac{9}{14} \log_2 \frac{9}{14} - \frac{5}{14} \log_2 \frac{5}{14} = 0.940$$

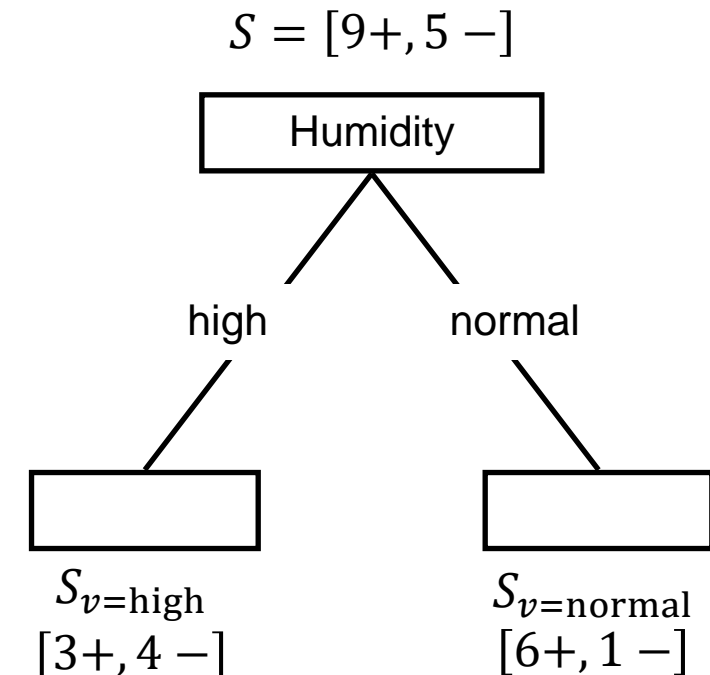
■ Step 2: Calculate entropy of child nodes

■ Entropy of $S_{v=\text{high}}$:

$$■ H(S_{v=\text{high}}) = -\frac{3}{7} \log_2 \frac{3}{7} - \frac{4}{7} \log_2 \frac{4}{7} = 0.985$$

■ Entropy of $S_{v=\text{normal}}$:

$$■ H(S_{v=\text{normal}}) = -\frac{6}{7} \log_2 \frac{6}{7} - \frac{1}{7} \log_2 \frac{1}{7} = 0.592$$



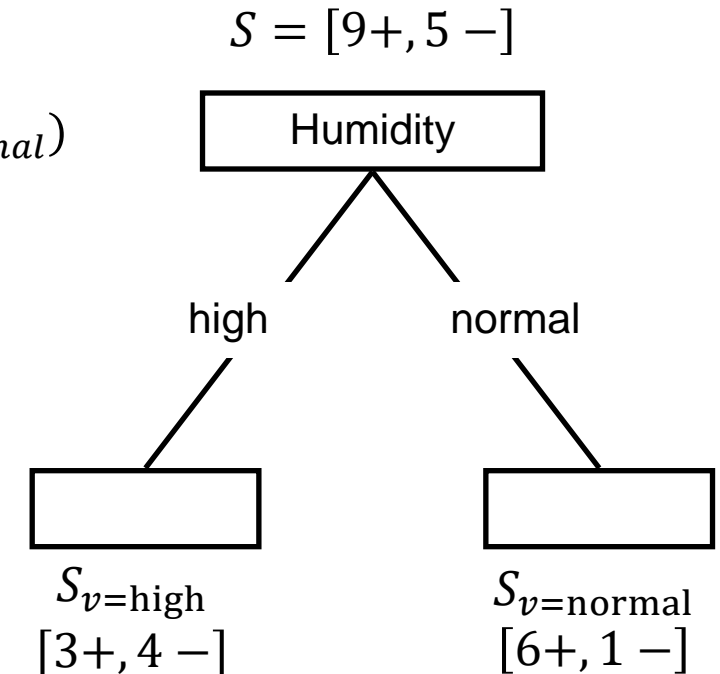
Information Gain – Attribute Selection

■ Step 3: Calculate Information gain for all attributes

■ Information gain for **humidity**:

$$■ \quad IG(S, Humidity) = H(S) - \frac{|S_{v=high}|}{|S|} H(S_{v=high}) - \frac{|S_{v=normal}|}{|S|} H(S_{v=normal})$$

$$■ \quad IG(S, Humidity) = 0.940 - \frac{7}{14} 0.985 - \frac{7}{14} 0.592 = 0.151$$



Information Gain – Attribute Selection

■ Step 3: Calculate Information gain for all attributes

■ Information gain for **humidity**:

$$■ IG(S, Humidity) = H(S) - \frac{|S_{v=high}|}{|S|} H(S_{v=high}) - \frac{|S_{v=normal}|}{|S|} H(S_{v=normal})$$

$$■ IG(S, Humidity) = 0.940 - \frac{7}{14} 0.985 - \frac{7}{14} 0.592 = 0.151$$

■ Information gain for **wind**:


$$■ H(S) = 0.940 \text{ (entropy of parent stays the same)}$$

$$■ H(S_{v=weak}) = -\frac{6}{8} \log_2 \frac{6}{8} - \frac{2}{8} \log_2 \frac{2}{8} = 0.811$$

$$■ H(S_{v=strong}) = -\frac{3}{6} \log_2 \frac{3}{6} - \frac{3}{6} \log_2 \frac{3}{6} = 1$$

$$■ IG(S, Wind) = H(S) - \frac{|S_{v=weak}|}{|S|} H(S_{v=weak}) - \frac{|S_{v=strong}|}{|S|} H(S_{v=strong})$$

$$■ IG(S, Wind) = 0.940 - \frac{8}{14} 0.811 - \frac{6}{14} 1 = 0.048$$



Information Gain – Attribute Selection

■ Step 3: Calculate information gain for all attributes

■ Information gain for **humidity**:

■ $IG(S, Humidity) = 0.151$


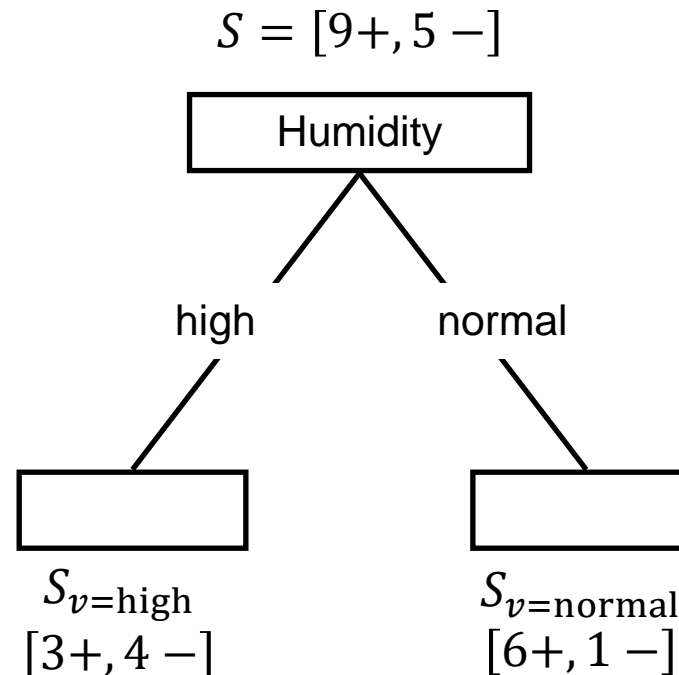
■ Information gain for **wind**:

■ $IG(S, Wind) = 0.048$

■ Step 4: Choose Attribute with highest information gain

■ $IG(S, Humidity) > IG(S, Wind)$

- Humidity attribute minimizes entropy, and is therefore used to split the tree




Overview

- Motivation
- Formalization
- Attribute selection
- Build the Tree - ID3-algorithm
- Overfitting
- Bagging
- Random Forest
- Extensions

Base Method: ID3

- **ID3**: Iterative Dichotomizer 3 \approx iterative division of instances
- **Top-Down**: Build decision tree **recursively** from root node
- **Greedy**: Each time an attribute is selected, the attribute that maximizes the information gain is used to split the data



Base Method: ID3 – Algorithm

ID3(Examples, Target Attribute, Attributes)

Create node for tree

IF (all examples positive), **Return** (node with label= \oplus)

IF (all examples negative), **Return** (node with label= \ominus)

IF (Attribute= \emptyset), **Return** (node with label= most common target attribute of examples in node)

Calculate A = Attribute with largest information gain for examples

Assign node the attribute = A

FOR ALL attribute values v_i of A :

Create new branch with v_i

Examples(v_i) = Subset of examples containing attribute value v_i

IF (Examples(v_i)= \emptyset):

Add leaf node to branch with label = most common target attribute of examples in node

THEN:

Add subtree **ID3**(Examples(v_i), Target Attribute, Attributes $\setminus \{A\}$)

Return node

Break recursion

Create subtree
and sort examples
into nodes

Recursively build
decision tree


Base Method: ID3 – Greedy Algorithm

■ Class is $A_1 \text{ XOR } A_2$

- Attribute A_3 does not correlate with final class
- A perfect tree tests A_1 then A_2 to make the classification with depth of tree = 2

■ BUT: $IG(S, A_1)$ and $IG(S, A_2) = 0$

- $H(S) = -\frac{4}{6}\log_2\frac{4}{6} - \frac{2}{6}\log_2\frac{2}{6} = 0.92$
- $H(S_0) = -\frac{2}{3}\log_2\frac{2}{3} - \frac{1}{3}\log_2\frac{1}{3} = 0.92$
- $H(S_1) = -\frac{1}{3}\log_2\frac{1}{3} - \frac{2}{3}\log_2\frac{2}{3} = 0.92$
- $IG(S, A_1) = 0.92 - \frac{3}{6}0.92 - \frac{3}{6}0.92 = 0$



A_1	A_2	A_3	Class
1	1	0	No
1	0	1	Yes
0	0	0	No
0	1	1	Yes
1	0	0	Yes
0	1	0	Yes

■ Information gain for split with A_3 is 0.25 and therefore preferred.


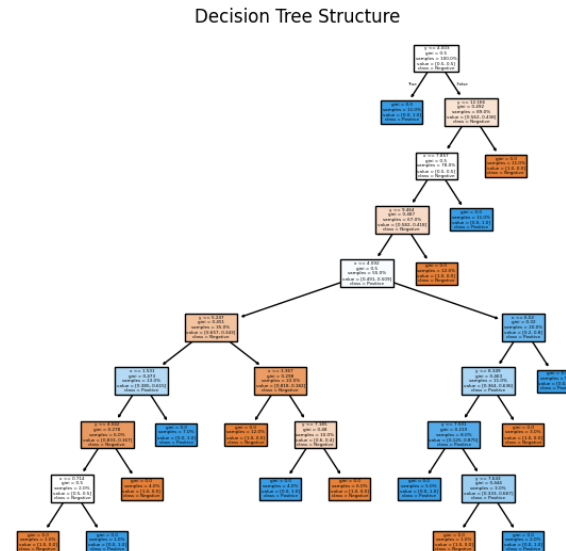

- \rightarrow ID3 choses A_3 and creates tree with depth = 3
- Greedy algorithms don't find the optimum. Finding an optimum tree is a NP-complete problem and is therefore rarely/never done. [\[Hyafil\]](#)

Base Method: ID3 – Properties

- The hypothesis space H of ID3 is the complete space of finite discrete-valued functions, relative to the available attributes → contains target function
- ID3 does not guarantee an optimal solution
 - Greedy algorithms do not backtrack and do not consider future steps
- Inductive bias: smaller trees are preferred to larger trees
- ID3 is prone to overfitting
 - Stops only when all instances have been classified perfectly, even if the data is noisy/wrong
- Nowadays modern libraries mostly use CART which is a slightly improved ID4.5 algorithm which is a slightly improved ID3 algorithm
 - The basics remain the same

Trees do not have an additive structure

- Decision Trees are unable to handle linearly correlated features.
- Example: Data from $y = x + 2$ and $y = x + 4$ with little bit of noise



Decision Tree is
unable to learn decision
boundary $y = x + 3$

*What is
happening
here?*

Overview

- Motivation
- Formalization
- Attribute selection
- Build the Tree - ID3-algorithm
- Overfitting
- Bagging
- Random Forest
- Extensions

Base Method ID3 – Overfitting

■ Building trees with ID3


- Add nodes until all training examples are perfectly classified
- Based on statistical approximation of information gain (e.g. entropy)

→ Often leads to overfitting because

- some examples in dataset are noisy or mislabeled
- examples are not representative (e.g. too little data)

Overfitting - Example

- **A, B:** Slightly different samples from same distribution
- Empirical error of train dataset is reduced to zero with deeper trees
- **BUT:**
 - Decision boundary on **A** seems to be suitable
 - The deeper decision boundaries on **B** include outliers and noisy samples. It overfits the data distribution



Occam's Razor

- Why should simpler hypotheses be preferred?
 - In this case, smaller trees?
- There are fewer simple hypotheses than complex ones.
 - A **short hypothesis** that correctly explains the training examples is most likely not a coincidence
 - A **long hypothesis**, that correctly explains the training examples may be a coincidence.
- Short trees are more **efficient**

Overfitting formal

- **Definition:** A hypothesis overfits the training examples, if some other hypothesis, that fits the training examples less well, actually performs better over the entire distribution of instances
 - Learning system memorizes training data rather than learning the underlying structure


- Formal:

$$h \in H \text{ overfitting} \Leftrightarrow \exists h' \in H \text{ such that given } D_{Tr} \text{ and } D_V \\ \hat{\mathcal{L}}_{D_{Tr}}(h) < \hat{\mathcal{L}}_{D_{Tr}}(h') \wedge \hat{\mathcal{L}}_{D_V}(h) > \hat{\mathcal{L}}_{D_V}(h')$$

- Whereas:
 - D_{Tr} – Training Data
 - D_V – Validation Data

Bias-Variance Tradeoff

$$\underbrace{\mathbb{E}_{x,y,D}[(h_D(x) - y)^2]}_{\text{Expected error}} = \underbrace{\mathbb{E}_{x,D}[(h_D(x) - \bar{h}(x))^2]}_{\text{Variance}} + \underbrace{\mathbb{E}_x[(\bar{h}(x) - \bar{y}(x))^2]}_{\text{Bias}^2} + \underbrace{\mathbb{E}_{x,y}[(\bar{y}(x) - y)^2]}_{\text{Noise}}$$




Common mistake in past exams: Forgetting the „squared“

Bias-Variance Tradeoff

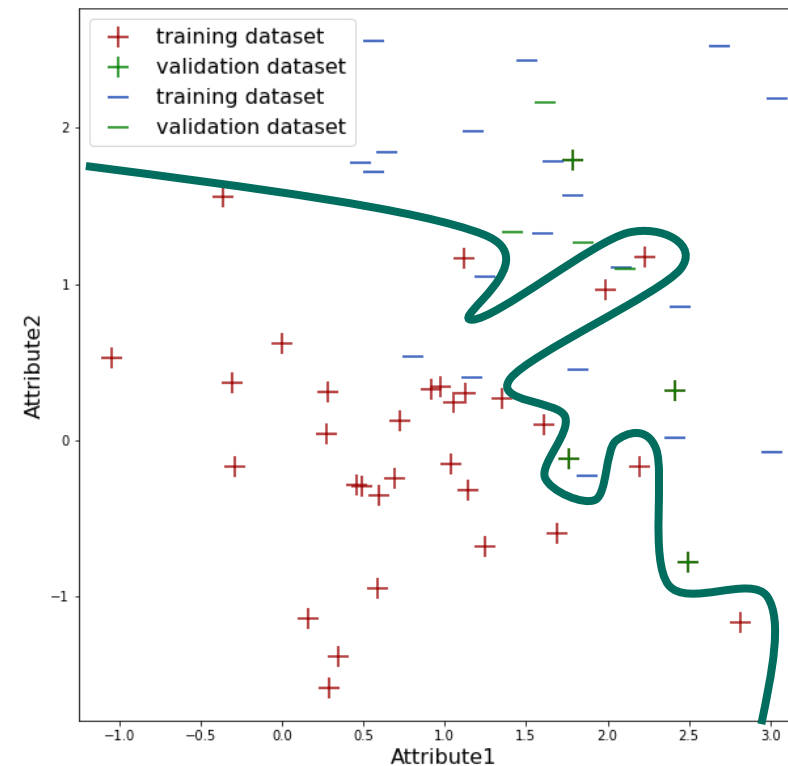
- **Variance:** Measures how much the classifier changes when trained on different splits of training data.
 - I.e. how much do we “overspecialize” for this particular training dataset
- **Bias:** What is the inherent error you get from your classifier even with infinite training data?
 - Can be related to the hypothesis space, e.g. a linear model can't predict non-linear data
- **Noise:** How large is the intrinsic noise in the data?
 - Measures ambiguity due to data distribution and feature representation.
 - Is an inherent aspect of the data and cannot be removed

Bias vs Variance

High bias, low variance
(usually underfitting)



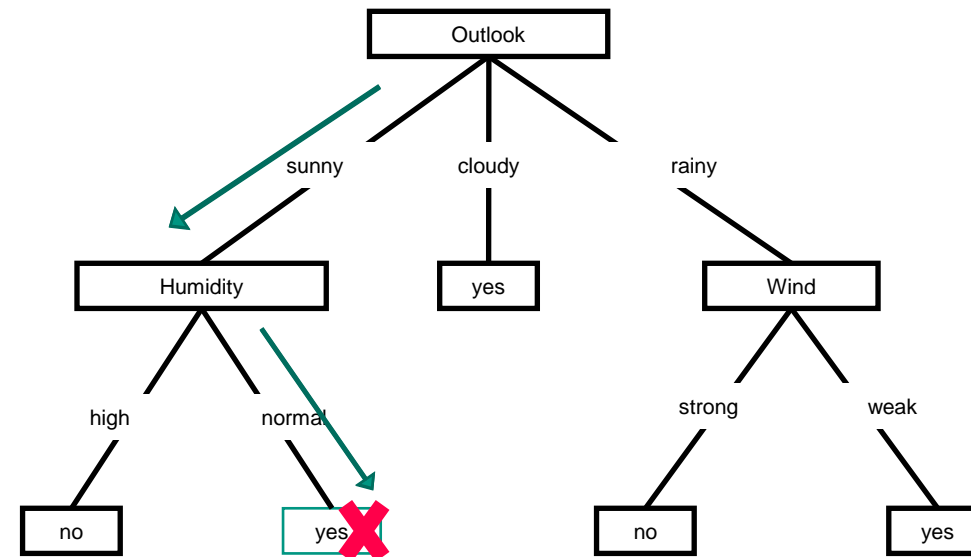
Low bias, high variance
(usually overfitting)



Example – Noisy/Wrong Data

- What happens if a noisy example is added to training data?

Nr.	Outlook	Temperature	Humidity	Wind	Tennis?
1	Sunny	Hot	Normal	Strong	No

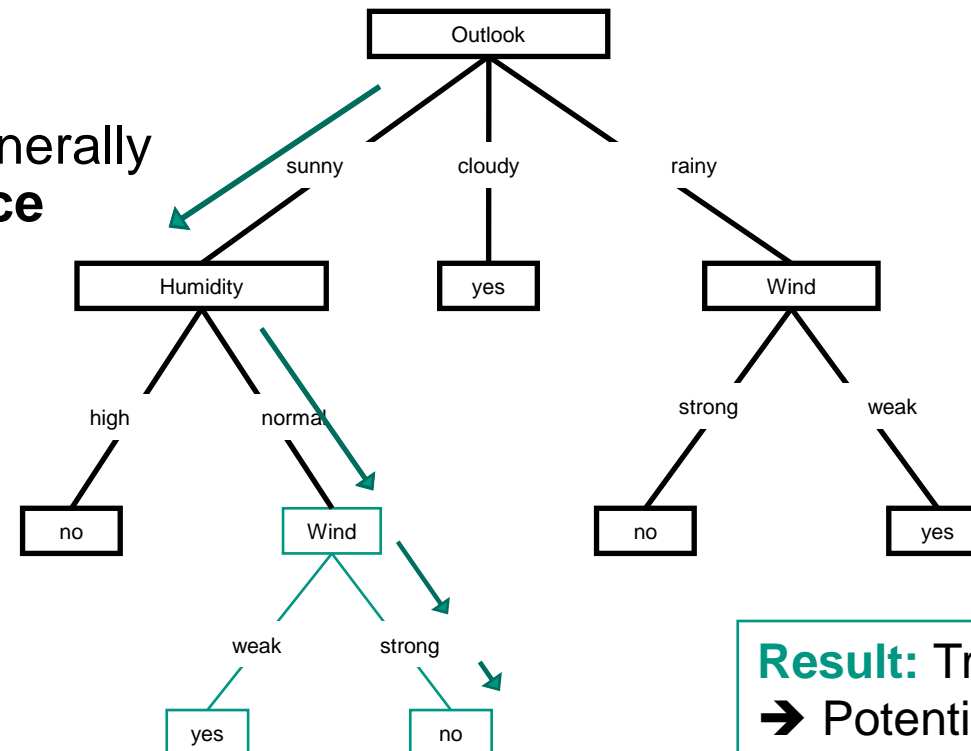


Example – Noisy/Wrong Data

- What happens if a noisy example is added to training data?

Nr.	Outlook	Temperature	Humidity	Wind	Tennis?
1	Sunny	Hot	Normal	Strong	No

- Decision Trees are generally **low bias high variance** models and usually **overfit** the data



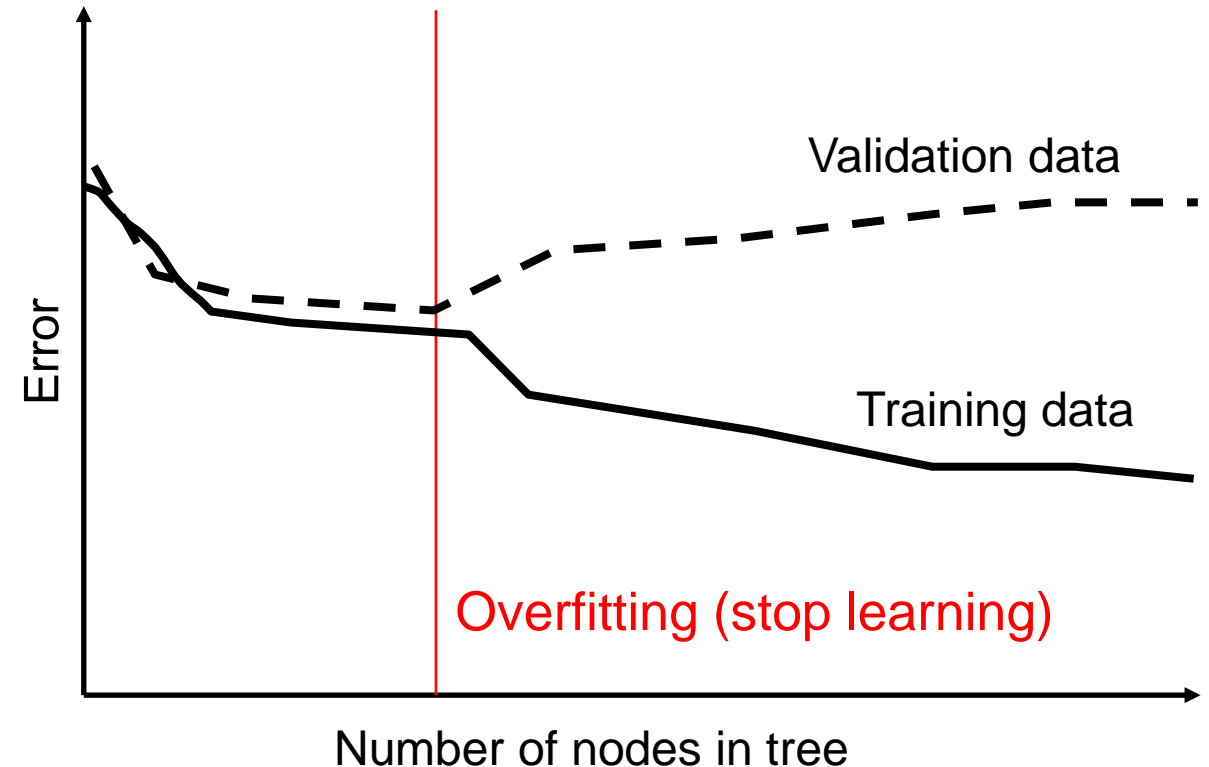
Result: Tree complexity increases
→ Potentially more errors on unseen data

Reduce Overfitting for ID3

- Combat the high overfitting properties of ID3 with **regularization**
 - **Maximum depth:** Growing the tree stops after a certain number of steps, even if not all data is perfectly classified.
 - **Minimum samples:** A node cannot be split if it contains less than a specified number of training examples
 - **Early stopping:** Stop growing the tree if the validation error increases
 - **Pruning:** Remove non-critical parts of the tree to reduce complexity
- Reduce overfitting with multiple trees:
 - **Bagging**
 - **Random Forests**


Early Stopping – Classification Error

- **Intuition:** Stop tree growth before overfitting occurs
- **Idea:** Validation error must decrease by ϵ per iteration, otherwise we stop
- **Pros:**
 - Easy to implement
- **Cons:**
 - Too short-sighted: Validation error could increase in the current step, but decrease again in the next step
 - Reminder the XOR problem with the greedy ID3 algorithm. Early stopping might stop before we even get the correct solution.




Pruning

- **Intuition:** Simplify the tree after learning is completed.
 - Stopping “too early” no longer a problem
- **Pruning:** Remove a node (subtree) and replace it with a leaf node of the most common class
- **Bottom up:** Starting from the leaves, replace each node with a leaf node until the validation error starts to increase



Pruning



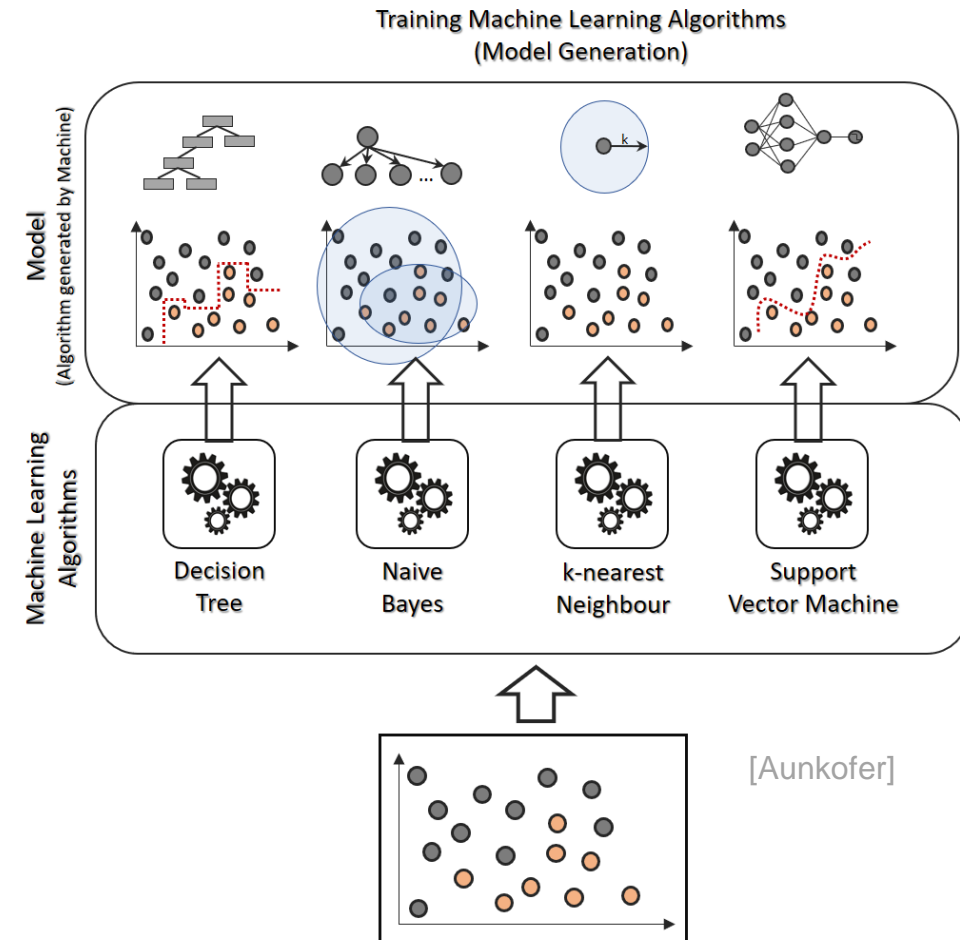
[Guestrin]

Overview

- Motivation
- Formalization
- Attribute selection
- Build the Tree - ID3-algorithm
- Overfitting
- Bagging
- Random Forest
- Extensions

Ensemble Learning

- **Question:** Which learning system is the best for my specific task?
 - Decision Trees
 - Support-Vector-Machines
 - Naive Bayes
 - ...?
- **Ensembles:** Learn several models and combine them into a stronger model



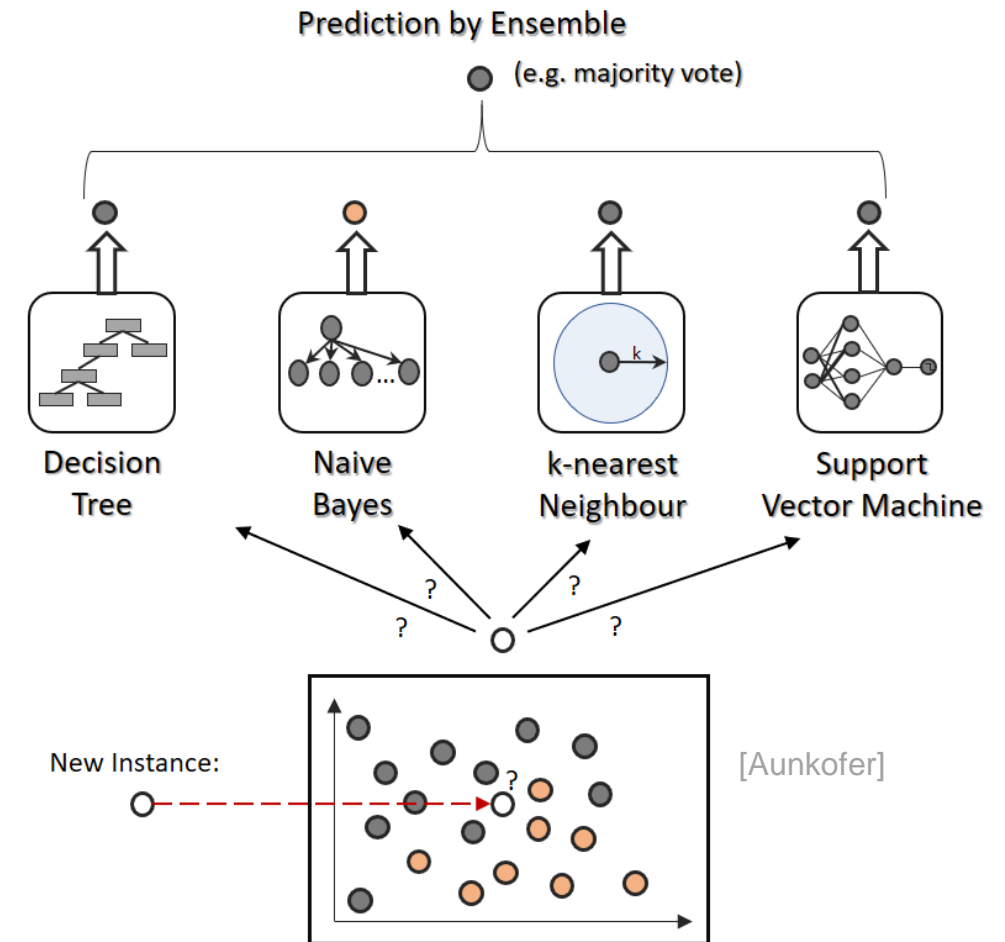
Ensemble Learning

■ Advantages

- Better prediction quality
- Better robustness (overfitting)

■ Disadvantages

- Time consuming and computationally expensive
- Loses interpretability
- Predictions from all machines can be treated the same or weighted differently
 - E.g.: use majority vote for classification or average for regression



Ensemble Learning - Motivation

- **Reminder:** Generalization Error = Bias² + Variance + Noise
- **Idea:** When several complex hypothesis with small bias and large variance are aggregated, their variance decreases but the bias remains small.
 - Simple way to decrease the generalization error


Ensemble Learning - Idealistic

- **Reminder:** Generalization Error = $\mathbb{E}_{x,D} \left[\left(h_D(x) - \bar{h}(x) \right)^2 \right] + \text{Bias}^2 + \text{Noise}$
- Instead of one dataset D , use several datasets D_1 to D_k
 - With the same number of learning examples and from the same distribution as D
- Learn k hypotheses with k different models
- Final hypothesis is e.g. the mean of the k hypotheses
- $\hat{h} = \frac{1}{k} \sum_{i=1}^k h_{D_i} \rightarrow \bar{h}$, if $k \rightarrow \infty$
 - Variance approaches 0!
- **Problem:** We don't have k datasets and if we split D into k datasets, the variance of the hypotheses increases even more due to the lack of training examples
- **Solution:** Bagging

Bagging

■ Process

- **Bootstrap:** Create k datasets D_1 to D_k from D via layback sampling
- **Training:** Train a model h_{D_1} to h_{D_k} for each dataset
- **Aggregation:** Combine models by majority vote or average their predictions
 - Regression: $h(x) = \frac{1}{k} \sum_{i=1}^k h_{D_i}(x)$
- **Advantage:** Reduce variance without increasing bias



Bagging – Pros & Cons

Pros

- **Reduces variance** and can therefore be used with high variance models
- **Uncertainty**: Bagging not only gives you the expected value, but also makes the variance of the prediction easily visible
- **Out-of-bag error**: Learning examples from D not sampled in dataset D_i can be used to validate h_i . No separate validation dataset is required.
- **Parallelizable**

Cons


- **Computational cost** increases by at least k
- Variance cannot be fully reduced because bootstrapping samples is not **I.I.D**
 - **I.I.D**: Independent and identically distributed random samples
- **Loss of interpretability**

Overview

- Motivation
- Formalization
- Attribute selection
- Build the Tree - ID3-algorithm
- Overfitting
- Bagging
- Random Forest
- Extensions

Random Forests



- **Problem of bagging:** Models are highly correlated
- **Previous:** Bagging with decision trees. Splits can use all d attributes of drawn samples
- **Now:** Choose random subset $s < d$ of attributes for each split. Create nodes using the best attribute for maximizing information gain in s .
- **Why:** To ensure, that the k learned trees are less correlated. Each tree uses a different subset of attributes.
 - Using s instead of d attributes will increase the bias. However, the reduction in variance from decorrelating the trees outweighs this.



Random Forests – Modification of Bagging

■ Process

- **Bootstrap:** Create k datasets D_1 to D_k from D via layback sampling.
- **Training:** Train a model h_{D_1} to h_{D_k} for each dataset with one small modification
 - **Before each split, sample a subset of $s < d$ attributes as possible candidates for splitting**
- **Aggregation:** Combine models by majority vote or average their predictions
- **Advantage:** Generally, much better results than single decision tree



Overview

- Motivation
- Formalization
- Attribute selection
- Build the Tree - ID3-algorithm
- Overfitting
- Bagging
- Random Forest
- Extensions

Continuous Attribute Values I

■ Given

- Attribute A with continuous values

■ Process

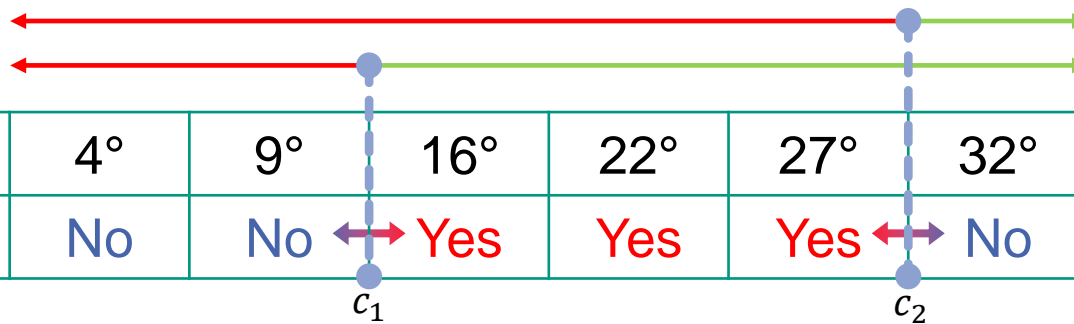
- Dynamic definition of a new discrete attribute: $A_c = \text{true}$ if $A > c$

■ Problem:

- Choice of threshold c ?
 - Use information gain:
 - Sort examples by their attribute values
 - Optimal threshold lies in the middle between two adjacent examples with different class affiliations

Continuous Attribute Values II

■ Example: Continuous temperature



Temperature	4°	9°	16°	22°	27°	32°
Tennis	No	No	Yes	Yes	Yes	No

c_1 c_2

■ Potential thresholds and respective information gain:

■ $c_1 = (9^\circ + 16^\circ) / 2 = 12.5^\circ \rightarrow IG = 1 - \frac{2}{6}0 - \frac{4}{6}0.81 = 0.46$

■ $c_2 = (27^\circ + 32^\circ) / 2 = 29.5^\circ \rightarrow IG = 1 - \frac{5}{6}0.97 - \frac{1}{6}0 = 0.19$

■ Information gain is higher $c_1 = 12.5^\circ$ categorical attributes are < 12.5 and > 12.5

Summary

- Decision trees are a **non-parametric** supervised learning method
- Generally, very **fast training and inference**
- **Interpretable** as an „if-else“ flow chart
- Decision Trees usually **overfit the data** and show **low accuracies** on test data
 - It's rare in practice to use a single decision tree
 - But ensemble methods like **random forests** and **boosting** drastically improve performance as they decrease variance
 - But we lose interpretability
 - Even today, they achieve better results on tabular data than neural networks
[\[2022: Why do tree-based models still outperform deep learning on tabular data\]](#)

Literature

- Tom Mitchell: *Machine Learning*, Chapter 3. 1997
 - Homepage: <http://www-2.cs.cmu.edu/~tom/>
 - Only individual Decision Trees, no ensembles
- Murphy: *Probabilistic Machine Learning*, Chapter 18. 2022
 - [PDF](#)
 - Contains ensemble methods like random forest and gradient boosting
 - Also includes XGBoost, which is currently one of the best tree-based methods.